



# Bits, Bytes, and Integers – Part 2

15-213/18-213/14-513/15-513/18-613: Introduction to Computer Systems  
3<sup>rd</sup> Lecture, September 3, 2019

# Assignment Announcements

## ■ Lab 0 available via course web page and Autolab.

- Due Thurs. Sept. 5, 11pm ET
- No grace days
- No late submissions
- Just do it!

## ■ Lab 1 available tonight via Autolab

- Due Thurs, Sept. 12, 11pm ET
- Read instructions carefully: writeup, bits.c, tests.c
  - Quirky software infrastructure
- Based on lectures 2, 3, and 4 (CS:APP Chapter 2)
- After today's lecture you will know everything for the integer problems
- Floating point covered Thursday Sept. 5

# Summary From Last Lecture

- **Representing information as bits**
- **Bit-level manipulations**
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
- Representations in memory, pointers, strings
- Summary


# Encoding Integers

## Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

## Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$



Sign Bit

## Two's Complement Examples ( $w = 5$ )

$$\begin{array}{rccccc} & -16 & 8 & 4 & 2 & 1 \\ 10 = & 0 & 1 & 0 & 1 & 0 & 8+2 = 10 \end{array}$$

$$\begin{array}{rccccc} & -16 & 8 & 4 & 2 & 1 \\ -10 = & 1 & 0 & 1 & 1 & 0 & -16+4+2 = -10 \end{array}$$

# Unsigned & Signed Numeric Values

$X$	$B2U(X)$	$B2T(X)$
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

## ■ Equivalence

- Same encodings for nonnegative values

## ■ Uniqueness


- Every bit pattern represents unique integer value
- Each representable integer has unique bit encoding

## ■ Expression containing signed and unsigned int:


`int` is cast to `unsigned`

# Sign Extension and Truncation

## ■ Sign Extension




## ■ Truncation



## ■ Misunderstanding integers can lead to the end of the world as we know it!

- Thule (Qaanaaq), Greenland
- US DoD “Site J” Ballistic Missile Early Warning System (BMEWS)
- 10/5/60: world nearly ends
- Missile radar echo: 1/8s
- BMEWS reports: 75s echo(!)
- 1000s of objects reported
- NORAD alert level 5:
  - Immediate incoming nuclear attack!!!!





- Kruschev was in NYC 10/5/60 (weird time to attack)
  - someone in Qaanaaq said “why not go check outside?”
- “Missiles” were actually THE MOON RISING OVER NORWAY
- Expected max distance: 3000 mi; Moon distance: .25M miles!
- .25M miles % sizeof(distance) = 2200mi.
- Overflow of distance nearly caused nuclear apocalypse!!

# Today: Bits, Bytes, and Integers


- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - **Addition, negation, multiplication, shifting**
- Representations in memory, pointers, strings
- Summary

# Unsigned Addition

Operands:  $w$  bits

True Sum:  $w+1$  bits

Discard Carry:  $w$  bits



## ■ Standard Addition Function

- Ignores carry output

## ■ Implements Modular Arithmetic

$$s = \text{UAdd}_w(u, v) = u + v \bmod 2^w$$

<code>unsigned char</code>	<code>1110 1001</code>	<code>E9</code>	<code>223</code>
	<code>+ 1101 0101</code>	<code>+ D5</code>	<code>+ 213</code>
	<code>1 1011 1110</code>	<code>1BE</code>	<code>446</code>
	<code>1011 1110</code>	<code>BE</code>	<code>190</code>


	Hex	Decimal	Binary
0	0	0000	
1	1	0001	
2	2	0010	
3	3	0011	
4	4	0100	
5	5	0101	
6	6	0110	
7	7	0111	
8	8	1000	
9	9	1001	
A	10	1010	
B	11	1011	
C	12	1100	
D	13	1101	
E	14	1110	
F	15	1111	

# Visualizing (Mathematical) Integer Addition

## ■ Integer Addition

- 4-bit integers  $u, v$
- Compute true sum  
 $\text{Add}_4(u, v)$
- Values increase linearly  
with  $u$  and  $v$
- Forms planar surface



$\text{Add}_4(u, v)$



# Visualizing Unsigned Addition

## Wraps Around

- If true sum  $\geq 2^w$
- At most once




# Two's Complement Addition

Operands:  $w$  bits

True Sum:  $w+1$  bits

Discard Carry:  $w$  bits



## ■ TAdd and UAdd have Identical Bit-Level Behavior

- Signed vs. unsigned addition in C:

```

int s, t, u, v;
s = (int) ((unsigned) u + (unsigned) v);
t = u + v
    
```

- Will give  $s == t$

$1110\ 1001$	$E9$	$-23$	
$+ 1101\ 0101$	$+ D5$	$+ -43$	
$\underline{\underline{1\ 1011\ 1110}}$		$\underline{\underline{1BE}}$	
$1011\ 1110$	$BE$	$-66$	


# TAdd Overflow

## ■ Functionality

- True sum requires  $w+1$  bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer

<b>0</b>	111...1
<b>0</b>	100...0
<b>0</b>	000...0
<b>1</b>	011...1
<b>1</b>	000...0

## True Sum



# Visualizing 2's Complement Addition

## ■ Values

- 4-bit two's comp.
- Range from -8 to +7

## ■ Wraps Around

- If  $\text{sum} \geq 2^{w-1}$ 
  - Becomes negative
  - At most once
- If  $\text{sum} < -2^{w-1}$ 
  - Becomes positive
  - At most once


NegOver

$TAdd_4(u, v)$

$u$

PosOver


$v$



# Characterizing TAdd

## ■ Functionality

- True sum requires  $w+1$  bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer




$$TAdd_w(u, v) = \begin{cases} u + v + 2^w & u + v < TMin_w \text{ (NegOver)} \\ u + v & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^w & TMax_w < u + v \text{ (PosOver)} \end{cases}$$

# Multiplication

- **Goal: Computing Product of  $w$ -bit numbers  $x, y$** 
  - Either signed or unsigned
- **But, exact results can be bigger than  $w$  bits**
  - Unsigned: up to  $2w$  bits
    - Result range:  $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
  - Two's complement min (negative): Up to  $2w-1$  bits
    - Result range:  $x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
  - Two's complement max (positive): Up to  $2w$  bits, but only for  $(TMin_w)^2$ 
    - Result range:  $x * y \leq (-2^{w-1})^2 = 2^{2w-2}$
- **So, maintaining exact results...**
  - would need to keep expanding word size with each product computed
  - is done in software, if needed
    - e.g., by “arbitrary precision” arithmetic packages

# Unsigned Multiplication in C

Operands:  $w$  bits



## ■ Standard Multiplication Function

- Ignores high order  $w$  bits


## ■ Implements Modular Arithmetic

$$\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$$

$1110\ 1001$	$E9$	$233$
$* \quad 1101\ 0101$	$* \quad D5$	$* \quad 213$
$1100\ 0001$	$C1DD$	$49629$
$1101\ 1101$	$DD$	$221$

# Signed Multiplication in C

Operands:  $w$  bits



## ■ Standard Multiplication Function

- Ignores high order  $w$  bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

$$\begin{array}{r}
 & 1110 & 1001 \\
 * & 1101 & 0101 \\
 \hline
 0000 & 0011 & 1101 & 1101 \\
 \hline
 1101 & 1101
 \end{array}
 \quad
 \begin{array}{r}
 E9 \\
 * D5 \\
 \hline
 03DD
 \end{array}
 \quad
 \begin{array}{r}
 -23 \\
 * -43 \\
 \hline
 989
 \end{array}
 \quad
 \begin{array}{r}
 DD \\
 \\
 -35
 \end{array}$$

# Power-of-2 Multiply with Shift


## ■ Operation

- $u \ll k$  gives  $u * 2^k$
- Both signed and unsigned

Operands:  $w$  bits

True Product:  $w+k$  bits

Discard  $k$  bits:  $w$  bits




## ■ Examples

- $u \ll 3 == u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$
- Most machines shift and add faster than multiply
  - Compiler generates this code automatically

# Unsigned Power-of-2 Divide with Shift

## ■ Quotient of Unsigned by Power of 2

- $u \gg k$  gives  $\lfloor u / 2^k \rfloor$
- Uses logical shift




	Division	Computed	Hex	Binary
x	15213	15213	3B 6D	00111011 01101101
x >> 1	7606.5	7606	1D B6	00011101 10110110
x >> 4	950.8125	950	03 B6	00000011 10110110
x >> 8	59.4257813	59	00 3B	00000000 00111011

# Signed Power-of-2 Divide with Shift

## ■ Quotient of Signed by Power of 2

- $x \gg k$  gives  $\lfloor x / 2^k \rfloor$
  - Uses arithmetic shift
  - Rounds wrong direction when  $x < 0$



	<b>Division</b>	<b>Computed</b>	<b>Hex</b>	<b>Binary</b>
x	-15213	-15213	C4 93	11000100 10010011
x >> 1	-7606.5	-7607	E2 49	11100010 01001001
x >> 4	-950.8125	-951	FC 49	11111100 01001001
x >> 8	-59.4257813	-60	FF C4	11111111 11000100


# Correct Power-of-2 Divide

## ■ Quotient of Negative Number by Power of 2

- Want  $\lceil x / 2^k \rceil$  (Round Toward 0)
- Compute as  $\lfloor (x+2^k-1) / 2^k \rfloor$ 
  - In C: `(x + (1<<k)-1) >> k`
  - Biases dividend toward 0

## Case 1: No rounding

Dividend:



*Biasing has no effect*

# Correct Power-of-2 Divide (Cont.)

## Case 2: Rounding

## Dividend:

The diagram illustrates the floating-point addition of  $x$  and  $+2^k - 1$ , resulting in a value incremented by 1. The binary representations are as follows:

- $x$ : A  $k$ -bit binary number starting with 1 (sign bit) followed by  $k-1$  bits.
- $+2^k - 1$ : A  $k$ -bit binary number starting with 0 followed by  $k-1$  bits.
- Result of Addition:** The sum of  $x$  and  $+2^k - 1$  is a  $k$ -bit binary number starting with 1 followed by  $k-1$  bits.
- Division by  $2^k$ :** The result is divided by  $2^k$  to obtain the normalized result  $x / 2^k$ .
- Normalized Result ( $x / 2^k$ ):** The result of the division is a  $k$ -bit binary number starting with 1 followed by  $k-1$  bits.
- Binary Point:** The binary point is indicated by a vertical line between the sign bit and the fraction bits.
- Incremented by 1:** The final result is the normalized result plus 1, represented by a  $k$ -bit binary number starting with 1 followed by  $k-1$  bits.

## ***Biasing adds 1 to final result***

# Negation: Complement & Increment

- Negate through complement and increase

$$\sim x + 1 == -x$$

- Example

- Observation:  $\sim x + x == 1111\dots111 == -1$

$$\begin{array}{r}
 x \quad \boxed{1} \boxed{0} \boxed{0} \boxed{1} \boxed{1} \boxed{1} \boxed{0} \boxed{1} \\
 + \quad \sim x \quad \boxed{0} \boxed{1} \boxed{1} \boxed{0} \boxed{0} \boxed{0} \boxed{1} \boxed{0} \\
 \hline
 -1 \quad \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1}
 \end{array}$$

$x = 15213$

	Decimal	Hex	Binary
$x$	15213	3B 6D	00111011 01101101
$\sim x$	-15214	C4 92	11000100 10010010
$\sim x + 1$	-15213	C4 93	11000100 10010011
$y$	-15213	C4 93	11000100 10010011

# Complement & Increment Examples

$x = 0$

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
$\sim 0$	-1	FF FF	11111111 11111111
$\sim 0+1$	0	00 00	00000000 00000000

$x = TMin$

	Decimal	Hex	Binary
x	-32768	80 00	10000000 00000000
$\sim x$	32767	7F FF	01111111 11111111
$\sim x+1$	-32768	80 00	10000000 00000000

## Canonical counter example

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - **Summary**
- Representations in memory, pointers, strings

# Arithmetic: Basic Rules

## ■ Addition:

- Unsigned/signed: Normal addition followed by truncate, same operation on bit level
- Unsigned: addition mod  $2^w$ 
  - Mathematical addition + possible subtraction of  $2^w$
- Signed: modified addition mod  $2^w$  (result in proper range)
  - Mathematical addition + possible addition or subtraction of  $2^w$

## ■ Multiplication:

- Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
- Unsigned: multiplication mod  $2^w$
- Signed: modified multiplication mod  $2^w$  (result in proper range)

# Quiz Time!


Check out:

<https://canvas.cmu.edu/courses/10968>

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- Representations in memory, pointers, strings

# Byte-Oriented Memory Organization



## ■ Programs refer to data by address

- Conceptually, envision it as a very large array of bytes
  - In reality, it's not, but can think of it that way
- An address is like an index into that array
  - and, a pointer variable stores an address

## ■ Note: system provides private address spaces to each “process”

- Think of a process as a program being executed
- So, a program can clobber its own data, but not that of others


# Machine Words

- Any given computer has a “Word Size”
  - Nominal size of integer-valued data
    - and of addresses
  - Until recently, most machines used 32 bits (4 bytes) as word size
    - Limits addresses to 4GB ( $2^{32}$  bytes)
  - Increasingly, machines have 64-bit word size
    - Potentially, could have 18 EB (exabytes) of addressable memory
    - That's  $18.4 \times 10^{18}$
  - Machines still support multiple data formats
    - Fractions or multiples of word size
    - Always integral number of bytes

# Word-Oriented Memory Organization

## ■ Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



# Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
<b>char</b>	1	1	1
<b>short</b>	2	2	2
<b>int</b>	4	4	4
<b>long</b>	4	8	8
<b>float</b>	4	4	4
<b>double</b>	8	8	8
<b>pointer</b>	4	8	8

# Byte Ordering


- So, how are the bytes within a multi-byte word ordered in memory?
- Conventions
  - Big Endian: Sun (Oracle SPARC), PPC Mac, *Internet*
    - Least significant byte has highest address
  - Little Endian: *x86*, ARM processors running Android, iOS, and Linux
    - Least significant byte has lowest address

# Byte Ordering Example


## ■ Example

- Variable x has 4-byte value of 0x01234567
- Address given by &x is 0x100

### BigEndian




### LittleEndian




# Representing Integers

```
int A = 15213;
```


Increasing addresses ↓



```
long int C = 15213;
```



```
int B = -15213;
```



**Two's complement representation**

# Examining Data Representations

## ■ Code to Print Byte Representation of Data

- Casting pointer to unsigned char \* allows treatment as a byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, size_t len){
    size_t i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}
```

### Printf directives:

%p: Print pointer  
%x: Print Hexadecimal

# show\_bytes Execution Example


```
int a = 15213;
printf("int a = 15213;\n");
show_bytes((pointer) &a, sizeof(int));
```

## Result (Linux x86-64):

```
int a = 15213;
0x7ffffb7f71dbc      6d
0x7ffffb7f71dbd      3b
0x7ffffb7f71dbe      00
0x7ffffb7f71dbf      00
```

# Representing Pointers

```
int B = -15213;  
int *P = &B;
```



Different compilers & machines assign different locations to objects

Even get different results each time run program

# Representing Strings


```
char S[6] = "18213";
```

## ■ Strings in C

- Represented by array of characters
- Each character encoded in ASCII format
  - Standard 7-bit encoding of character set
  - Character “0” has code 0x30
    - Digit  $i$  has code  $0x30+i$
- String should be null-terminated
  - Final character = 0

## ■ Compatibility

- Byte ordering not an issue



# Reading Byte-Reversed Listings

## ■ Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

## ■ Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

## ■ Deciphering Numbers

- Value:
- Pad to 32 bits:
- Split into bytes:
- Reverse:

0x12ab

0x000012ab

00 00 12 ab

ab 12 00 00

# Summary

- **Representing information as bits**
- **Bit-level manipulations**
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
- **Representations in memory, pointers, strings**
- **Summary**